

Imports

```
1 import os
2 import warnings
3 import numpy as np
4 import pandas as pd
5 import matplotlib.pyplot as plt
6 from PIL import Image
7
8 import torch
9 import torch.nn as nn
10 import torch.optim as optim
11 from torch.utils.data import Dataset, DataLoader
12 from torchvision import transforms
13
14 from sklearn.model_selection import train_test_split
15 from sklearn.preprocessing import StandardScaler
16 from tqdm import tqdm
17
```

```
1 warnings.filterwarnings("ignore")
```

Configuration & Reproducibility

```
1 SEED = 42
2 BASE_PATH = "/kaggle/input/csiro-biomass"
3
4 IMG_SIZE = 224
5 BATCH_SIZE = 16
6 NUM_EPOCHS = 50
7
8 LEARNING_RATE = 1e-3
9 WEIGHT_DECAY = 1e-4
10
11 DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
12
13 np.random.seed(SEED)
14 torch.manual_seed(SEED)
15 torch.cuda.manual_seed_all(SEED)
16 torch.backends.cudnn.deterministic = True
17 torch.backends.cudnn.benchmark = False
18
```

Load & Reshape Training Data

```
1 train_csv = pd.read_csv(f"{BASE_PATH}/train.csv")
2
3 train_df = train_csv.pivot_table(
4     index="image_path",
5     columns="target_name",
6     values="target"
7 ).reset_index()
8
9 train_df.fillna(0, inplace=True)
10
```

```
1 TARGET_COLUMNS = [
2     'Dry_Green_g',
3     'Dry_Dead_g',
4     'Dry_Clover_g',
5     'GDM_g',
6     'Dry_Total_g'
```

7]
8

▼ Train / Validation Split & Target Scaling

```
1 train_df, val_df = train_test_split(  
2     train_df, test_size=0.2, random_state=SEED, shuffle=True  
3 )  
4
```

```
1 target_scaler = StandardScaler()  
2 target_scaler.fit(train_df[TARGET_COLUMNS].values)  
3
```

▼ **StandardScaler** ⓘ ⓘ
StandardScaler()

▼ Dataset Definition

```
1 class BiomassDataset(Dataset):  
2     def __init__(self, dataframe, image_dir, transform, scaler=None):  
3         self.df = dataframe.copy()  
4         self.df["image_path"] = self.df["image_path"].apply(os.path.basename)  
5         self.samples = list(self.df.itertuples(index=False))  
6         self.image_dir = image_dir  
7         self.transform = transform  
8         self.scaler = scaler  
9  
10    def __len__(self):  
11        return len(self.samples)  
12  
13    def __getitem__(self, idx):  
14        row = self.samples[idx]  
15  
16        image = Image.open(  
17            os.path.join(self.image_dir, row.image_path)  
18        ).convert("RGB")  
19        image = self.transform(image)  
20  
21        targets = np.array(  
22            [getattr(row, c) for c in TARGET_COLUMNS],  
23            dtype=np.float32  
24        )  
25  
26        if self.scaler:  
27            targets = self.scaler.transform(targets.reshape(1, -1))[0]  
28  
29        return image, torch.tensor(targets)  
30
```

▼ Transforms & DataLoaders

```
1 image_transforms = transforms.Compose([  
2     transforms.Resize((IMG_SIZE, IMG_SIZE)),  
3     transforms.ToTensor(),  
4     transforms.Normalize(  
5         mean=[0.485, 0.456, 0.406],  
6         std=[0.229, 0.224, 0.225]  
7     )  
8 ])  
9
```

```

1 train_loader = DataLoader(
2     BiomassDataset(train_df, f"{BASE_PATH}/train", image_transforms, target_scaler),
3     batch_size=BATCH_SIZE, shuffle=True, num_workers=2
4 )
5
6 val_loader = DataLoader(
7     BiomassDataset(val_df, f"{BASE_PATH}/train", image_transforms, target_scaler),
8     batch_size=BATCH_SIZE, shuffle=False, num_workers=2
9 )
10

```

▼ CNN Regression Model

```

1 class BiomassCNN(nn.Module):
2     def __init__(self):
3         super().__init__()
4
5         self.features = nn.Sequential(
6             nn.Conv2d(3, 32, 3, padding=1),
7             nn.ReLU(),
8             nn.MaxPool2d(2),
9
10            nn.Conv2d(32, 64, 3, padding=1),
11            nn.ReLU(),
12            nn.MaxPool2d(2),
13
14            nn.Conv2d(64, 128, 3, padding=1),
15            nn.ReLU(),
16            nn.AdaptiveAvgPool2d(1)
17        )
18
19        self.regressor = nn.Linear(128, len(TARGET_COLUMNS))
20
21    def forward(self, x):
22        x = self.features(x)
23        return self.regressor(x.flatten(1))
24

```

```

1 model = BiomassCNN().to(DEVICE)
2

```

▼ Loss Function, Optimizer & Metric

```

1 TARGET_WEIGHTS = torch.tensor(
2     [0.1, 0.1, 0.1, 0.2, 0.5],
3     device=DEVICE
4 )
5

```

```

1 class WeightedMSELoss(nn.Module):
2     def __init__(self, weights):
3         super().__init__()
4         self.register_buffer("weights", weights)
5
6     def forward(self, predictions, targets):
7         return ((predictions - targets) ** 2 * self.weights).mean()
8

```

```

1 criterion = WeightedMSELoss(TARGET_WEIGHTS)
2 optimizer = optim.Adam(
3     model.parameters(),
4     lr=LEARNING_RATE,
5     weight_decay=WEIGHT_DECAY
6 )
7

```

```

1 def weighted_r2_score(y_true, y_pred, weights):
2     mean_y = y_true.mean(0, keepdim=True)
3     ss_res = ((y_true - y_pred) ** 2).sum(0)
4     ss_tot = ((y_true - mean_y) ** 2).sum(0)
5     r2 = 1 - (weights * ss_res).sum() / ((weights * ss_tot).sum() + 1e-8)
6     return r2.item()
7

```

✓ Training / Validation Epoch Function

```

1 def run_epoch(model, dataloader, optimizer=None):
2     training = optimizer is not None
3     model.train() if training else model.eval()
4
5     total_loss, total_r2, count = 0.0, 0.0, 0
6
7     with torch.set_grad_enabled(training):
8         for images, targets in dataloader:
9             images = images.to(DEVICE)
10            targets = targets.to(DEVICE)
11
12            predictions = model(images)
13            loss = criterion(predictions, targets)
14
15            if training:
16                optimizer.zero_grad()
17                loss.backward()
18                optimizer.step()
19
20            batch_size = images.size(0)
21            total_loss += loss.item() * batch_size
22            total_r2 += weighted_r2_score(targets, predictions, TARGET_WEIGHTS) * batch_size
23            count += batch_size
24
25    return total_loss / count, total_r2 / count
26

```

✓ Training Loop with History Tracking

```

1 train_losses, val_losses = [], []
2 train_r2s, val_r2s = [], []
3
4 best_val_r2 = -1e9
5
6 for epoch in range(NUM_EPOCHS):
7     train_loss, train_r2 = run_epoch(model, train_loader, optimizer)
8     val_loss, val_r2 = run_epoch(model, val_loader)
9
10    train_losses.append(train_loss)
11    val_losses.append(val_loss)
12    train_r2s.append(train_r2)
13    val_r2s.append(val_r2)
14
15    if val_r2 > best_val_r2:
16        best_val_r2 = val_r2
17        torch.save(model.state_dict(), "best_model.pth")
18
19    print(
20        f"Epoch {epoch+1}/{NUM_EPOCHS} | "
21        f"Train Loss: {train_loss:.4f}, R2: {train_r2:.4f} | "
22        f"Val Loss: {val_loss:.4f}, R2: {val_r2:.4f}"
23    )
24

```

```

Epoch 1/50 | Train Loss: 0.2013, R2: -0.1382 | Val Loss: 0.1561, R2: -0.0901
Epoch 2/50 | Train Loss: 0.1956, R2: -0.0285 | Val Loss: 0.1466, R2: -0.0304
Epoch 3/50 | Train Loss: 0.1838, R2: 0.0347 | Val Loss: 0.1373, R2: 0.0016
Epoch 4/50 | Train Loss: 0.1781, R2: 0.0654 | Val Loss: 0.1318, R2: 0.0451
Epoch 5/50 | Train Loss: 0.1736, R2: 0.0968 | Val Loss: 0.1295, R2: 0.0782

```

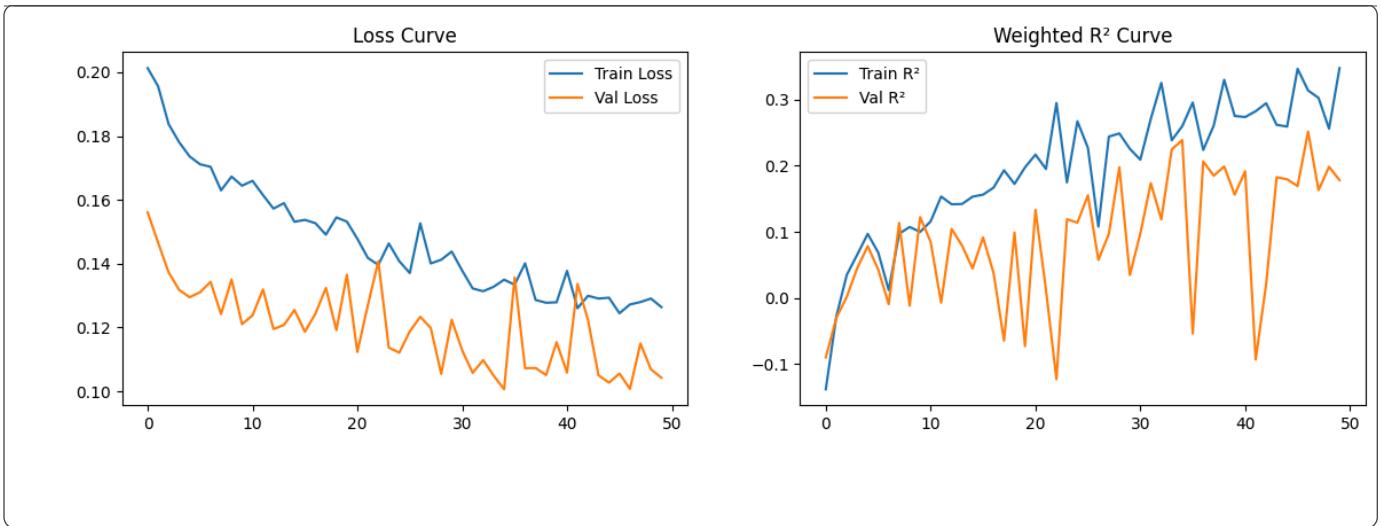
Epoch 6/50	Train Loss: 0.1712, R ² : 0.0686	Val Loss: 0.1311, R ² : 0.0428
Epoch 7/50	Train Loss: 0.1704, R ² : 0.0118	Val Loss: 0.1343, R ² : -0.0096
Epoch 8/50	Train Loss: 0.1630, R ² : 0.0971	Val Loss: 0.1242, R ² : 0.1134
Epoch 9/50	Train Loss: 0.1673, R ² : 0.1071	Val Loss: 0.1351, R ² : -0.0118
Epoch 10/50	Train Loss: 0.1645, R ² : 0.0998	Val Loss: 0.1211, R ² : 0.1220
Epoch 11/50	Train Loss: 0.1660, R ² : 0.1156	Val Loss: 0.1238, R ² : 0.0847
Epoch 12/50	Train Loss: 0.1615, R ² : 0.1533	Val Loss: 0.1319, R ² : -0.0073
Epoch 13/50	Train Loss: 0.1573, R ² : 0.1415	Val Loss: 0.1195, R ² : 0.1043
Epoch 14/50	Train Loss: 0.1590, R ² : 0.1420	Val Loss: 0.1208, R ² : 0.0789
Epoch 15/50	Train Loss: 0.1531, R ² : 0.1531	Val Loss: 0.1255, R ² : 0.0444
Epoch 16/50	Train Loss: 0.1538, R ² : 0.1561	Val Loss: 0.1187, R ² : 0.0914
Epoch 17/50	Train Loss: 0.1527, R ² : 0.1667	Val Loss: 0.1243, R ² : 0.0379
Epoch 18/50	Train Loss: 0.1491, R ² : 0.1930	Val Loss: 0.1324, R ² : -0.0647
Epoch 19/50	Train Loss: 0.1545, R ² : 0.1723	Val Loss: 0.1192, R ² : 0.0988
Epoch 20/50	Train Loss: 0.1532, R ² : 0.1973	Val Loss: 0.1365, R ² : -0.0730
Epoch 21/50	Train Loss: 0.1478, R ² : 0.2167	Val Loss: 0.1123, R ² : 0.1331
Epoch 22/50	Train Loss: 0.1418, R ² : 0.1948	Val Loss: 0.1270, R ² : 0.0123
Epoch 23/50	Train Loss: 0.1396, R ² : 0.2947	Val Loss: 0.1408, R ² : -0.1231
Epoch 24/50	Train Loss: 0.1464, R ² : 0.1746	Val Loss: 0.1137, R ² : 0.1189
Epoch 25/50	Train Loss: 0.1408, R ² : 0.2670	Val Loss: 0.1121, R ² : 0.1137
Epoch 26/50	Train Loss: 0.1371, R ² : 0.2269	Val Loss: 0.1188, R ² : 0.1551
Epoch 27/50	Train Loss: 0.1527, R ² : 0.1078	Val Loss: 0.1234, R ² : 0.0574
Epoch 28/50	Train Loss: 0.1401, R ² : 0.2439	Val Loss: 0.1199, R ² : 0.0968
Epoch 29/50	Train Loss: 0.1412, R ² : 0.2487	Val Loss: 0.1055, R ² : 0.1974
Epoch 30/50	Train Loss: 0.1438, R ² : 0.2254	Val Loss: 0.1224, R ² : 0.0345
Epoch 31/50	Train Loss: 0.1378, R ² : 0.2090	Val Loss: 0.1129, R ² : 0.0974
Epoch 32/50	Train Loss: 0.1322, R ² : 0.2709	Val Loss: 0.1058, R ² : 0.1735
Epoch 33/50	Train Loss: 0.1314, R ² : 0.3252	Val Loss: 0.1098, R ² : 0.1186
Epoch 34/50	Train Loss: 0.1328, R ² : 0.2383	Val Loss: 0.1049, R ² : 0.2246
Epoch 35/50	Train Loss: 0.1350, R ² : 0.2597	Val Loss: 0.1007, R ² : 0.2385
Epoch 36/50	Train Loss: 0.1334, R ² : 0.2955	Val Loss: 0.1357, R ² : -0.0540
Epoch 37/50	Train Loss: 0.1401, R ² : 0.2237	Val Loss: 0.1073, R ² : 0.2063
Epoch 38/50	Train Loss: 0.1286, R ² : 0.2605	Val Loss: 0.1073, R ² : 0.1848
Epoch 39/50	Train Loss: 0.1278, R ² : 0.3298	Val Loss: 0.1051, R ² : 0.1987
Epoch 40/50	Train Loss: 0.1279, R ² : 0.2751	Val Loss: 0.1154, R ² : 0.1561
Epoch 41/50	Train Loss: 0.1378, R ² : 0.2734	Val Loss: 0.1059, R ² : 0.1915
Epoch 42/50	Train Loss: 0.1261, R ² : 0.2823	Val Loss: 0.1337, R ² : -0.0932
Epoch 43/50	Train Loss: 0.1300, R ² : 0.2943	Val Loss: 0.1224, R ² : 0.0208
Epoch 44/50	Train Loss: 0.1291, R ² : 0.2617	Val Loss: 0.1050, R ² : 0.1824
Epoch 45/50	Train Loss: 0.1293, R ² : 0.2590	Val Loss: 0.1027, R ² : 0.1793
Epoch 46/50	Train Loss: 0.1245, R ² : 0.3464	Val Loss: 0.1056, R ² : 0.1691
Epoch 47/50	Train Loss: 0.1272, R ² : 0.3136	Val Loss: 0.1008, R ² : 0.2515
Epoch 48/50	Train Loss: 0.1280, R ² : 0.3024	Val Loss: 0.1150, R ² : 0.1629
Epoch 49/50	Train Loss: 0.1291, R ² : 0.2558	Val Loss: 0.1069, R ² : 0.1982
Epoch 50/50	Train Loss: 0.1264, R ² : 0.3475	Val Loss: 0.1043, R ² : 0.1780

▼ Training Curves (Loss & R²)

```

1 plt.figure(figsize=(14,4))
2
3 plt.subplot(1,2,1)
4 plt.plot(train_losses, label="Train Loss")
5 plt.plot(val_losses, label="Val Loss")
6 plt.legend()
7 plt.title("Loss Curve")
8
9 plt.subplot(1,2,2)
10 plt.plot(train_r2s, label="Train R2")
11 plt.plot(val_r2s, label="Val R2")
12 plt.legend()
13 plt.title("Weighted R2 Curve")
14
15 plt.show()
16

```



▼ Load Best Model

```

1 best_model = BiomassCNN().to(DEVICE)
2 best_model.load_state_dict(torch.load("best_model.pth", map_location=DEVICE))
3 best_model.eval()
4

BiomassCNN(
    (features): Sequential(
        (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (1): ReLU()
        (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (4): ReLU()
        (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (6): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (7): ReLU()
        (8): AdaptiveAvgPool2d(output_size=1)
    )
    (regressor): Linear(in_features=128, out_features=5, bias=True)
)

```

▼ Test Inference & Submission

```

1 test_dir = f"{BASE_PATH}/test"
2 submission_rows = []
3
4 with torch.no_grad():
5     for file in tqdm(sorted(os.listdir(test_dir))):
6         image = Image.open(os.path.join(test_dir, file)).convert("RGB")
7         tensor = image_transforms(image).unsqueeze(0).to(DEVICE)
8
9         preds_scaled = best_model(tensor).cpu().numpy()
10        preds = target_scaler.inverse_transform(preds_scaled)[0]
11
12        image_id = file.replace(".jpg", "")
13        for name, value in zip(TARGET_COLUMNS, preds):
14            submission_rows.append({
15                "sample_id": f"{image_id}__{name}",
16                "target": max(0, value)
17            })
18

```

100% |██████████| 1/1 [00:00<00:00, 6.92it/s]

```

1 submission_df = pd.DataFrame(submission_rows)
2 submission_df.head()
3

```

	sample_id	target
0	ID1001187975__Dry_Green_g	17.660715
1	ID1001187975__Dry_Dead_g	18.764156
2	ID1001187975__Dry_Clover_g	2.354568
3	ID1001187975__GDM_g	20.842461
4	ID1001187975__Dry_Total_g	38.504654

```
1 submission_df.to_csv("submission.csv", index=False)
```