

## Imports

```
1 import os
2 import warnings
3 import numpy as np
4 import pandas as pd
5 import matplotlib.pyplot as plt
6 from PIL import Image
7
8 import torch
9 import torch.nn as nn
10 import torch.optim as optim
11 from torch.utils.data import Dataset, DataLoader
12 from torchvision import transforms
13
14 from sklearn.model_selection import train_test_split
15 from sklearn.preprocessing import StandardScaler
16 from tqdm import tqdm
17 import joblib
18
19 warnings.filterwarnings("ignore")
20
```

## Config & Reproducibility

```
1 SEED = 42
2 BASE_PATH = "/kaggle/input/csiro-biomass"
3 IMG_SIZE = 224
4 BATCH_SIZE = 16
5 NUM_EPOCHS = 50
6 LEARNING_RATE = 1e-3
7 WEIGHT_DECAY = 1e-4
8
9 DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
10
11 np.random.seed(SEED)
12 torch.manual_seed(SEED)
13 torch.cuda.manual_seed_all(SEED)
14 torch.backends.cudnn.deterministic = True
15 torch.backends.cudnn.benchmark = False
16
```

## Load & Pivot Dataset

```

1 df = pd.read_csv(os.path.join(BASE_PATH, "train.csv"))
2
3 df = df.pivot_table(
4     index="image_path",
5     columns="target_name",
6     values="target"
7 ).reset_index()
8
9 df.fillna(0, inplace=True)
10
11 TARGET_COLUMNS = [
12     'Dry_Green_g',
13     'Dry_Dead_g',
14     'Dry_Clover_g',
15     'GDM_g',
16     'Dry_Total_g'
17 ]
18
19 print("Dataset shape:", df.shape)
20

```

Dataset shape: (357, 6)

## ▼ Train / Validation Split & Scaling

```

1 train_df, val_df = train_test_split(
2     df,
3     test_size=0.2,
4     random_state=SEED,
5     shuffle=True
6 )
7
8 target_scaler = StandardScaler()
9 target_scaler.fit(train_df[TARGET_COLUMNS].values)
10

```

▼ **StandardScaler** ⓘ ⓘ  
StandardScaler()

## ▼ Dataset Class

```

1 class BiomassDataset(Dataset):
2     def __init__(self, df, image_dir, transform=None, scaler=None):
3         self.df = df.copy()
4         self.df["image_path"] = self.df["image_path"].apply(os.path.basename)
5         self.samples = list(self.df.itertuples(index=False))
6         self.image_dir = image_dir
7         self.transform = transform
8         self.scaler = scaler
9
10    def __len__(self):
11        return len(self.samples)
12
13    def __getitem__(self, idx):
14        row = self.samples[idx]
15        image = Image.open(
16            os.path.join(self.image_dir, row.image_path)
17        ).convert("RGB")
18
19        if self.transform:
20            image = self.transform(image)
21
22        target = np.array(
23            [getattr(row, t) for t in TARGET_COLUMNS],
24            dtype=np.float32
25        )
26

```

```

27     if self.scaler:
28         target = self.scaler.transform(target.reshape(1, -1))[0]
29
30     return image, torch.tensor(target)
31

```

## Transforms (Medium Augmentation)

```

1 train_transforms = transforms.Compose([
2     transforms.Resize((IMG_SIZE, IMG_SIZE)),
3     transforms.RandomHorizontalFlip(p=0.5),
4     transforms.RandomRotation(10),
5     transforms.ColorJitter(brightness=0.2, contrast=0.2),
6     transforms.ToTensor(),
7     transforms.Normalize([0.485,0.456,0.406],[0.229,0.224,0.225])
8 ])
9
10 val_transforms = transforms.Compose([
11     transforms.Resize((IMG_SIZE, IMG_SIZE)),
12     transforms.ToTensor(),
13     transforms.Normalize([0.485,0.456,0.406],[0.229,0.224,0.225])
14 ])
15

```

## DataLoaders

```

1 train_dataset = BiomassDataset(
2     train_df,
3     os.path.join(BASE_PATH, "train"),
4     train_transforms,
5     target_scaler
6 )
7
8 val_dataset = BiomassDataset(
9     val_df,
10    os.path.join(BASE_PATH, "train"),
11    val_transforms,
12    target_scaler
13 )
14
15 train_loader = DataLoader(
16     train_dataset,
17     batch_size=BATCH_SIZE,
18     shuffle=True,
19     num_workers=2
20 )
21
22 val_loader = DataLoader(
23     val_dataset,
24     batch_size=BATCH_SIZE,
25     shuffle=False,
26     num_workers=2
27 )
28

```

## Weighted Loss & Metric

```

1 TARGET_WEIGHTS = torch.tensor([0.1,0.1,0.1,0.2,0.5], device=DEVICE)
2
3 class WeightedMSELoss(nn.Module):
4     def __init__(self, weights):
5         super().__init__()
6         self.register_buffer("weights", weights)
7
8     def forward(self, preds, targets):

```

```

9         return ((preds - targets) ** 2 * self.weights).mean()
10
11 criterion = WeightedMSELoss(TARGET_WEIGHTS)
12
13 def weighted_r2_score(y_true, y_pred, weights):
14     mean_y = y_true.mean(0, keepdim=True)
15     ss_res = ((y_true - y_pred) ** 2).sum(0)
16     ss_tot = ((y_true - mean_y) ** 2).sum(0)
17     return (1 - (weights * ss_res).sum() / ((weights * ss_tot).sum() + 1e-8)).item()
18

```

## Optimized CNN Architecture

```

1 class OptimizedCNN(nn.Module):
2     def __init__(self, num_outputs=5):
3         super().__init__()
4         layers = []
5         in_channels = 3
6         num_filters = 256
7
8         # EXPERIMENT WINNER: Kernel Size 3x3 (Best Val R2: 0.238753)
9         kernel_size = 3
10        padding = kernel_size // 2
11
12        # EXPERIMENT WINNER: Dropout 0.6 (Best Val R2: 0.260552 with wd 0.0)
13        dropout_p = 0.6
14
15        # EXPERIMENT WINNER: 3 Blocks (Strong performance, balancing depth and loss)
16        for _ in range(3):
17            layers += [
18                nn.Conv2d(in_channels, num_filters, kernel_size, padding=padding),
19                nn.BatchNorm2d(num_filters),
20                nn.ReLU(inplace=True),
21                nn.MaxPool2d(2),
22                nn.Dropout(dropout_p)
23            ]
24            in_channels = num_filters
25
26        layers.append(nn.AdaptiveAvgPool2d(1))
27        self.features = nn.Sequential(*layers)
28        self.regressor = nn.Linear(num_filters, num_outputs)
29
30    def forward(self, x):
31        x = self.features(x)
32        return self.regressor(x.flatten(1))
33
34 # --- TRAINING CONFIGURATION ---
35 # EXPERIMENT WINNER: Weight Decay 0.0 (per the dropout_0.6_wd_0.0 config)
36 model = OptimizedCNN(num_outputs=5)
37 optimizer = torch.optim.Adam(
38     model.parameters(),
39     lr=0.001,
40     weight_decay=0.0 # Set to 0.0 based on your best R2 result
41 )

```

## Epoch Function

```

1 def run_epoch(model, loader, optimizer=None):
2     training = optimizer is not None
3     model.train() if training else model.eval()
4
5     total_loss, total_r2, n = 0.0, 0.0, 0
6
7     with torch.set_grad_enabled(training):
8         for images, targets in loader:
9             images, targets = images.to(DEVICE), targets.to(DEVICE)
10            preds = model(images)
11            loss = criterion(preds, targets)

```

```

12
13     if training:
14         optimizer.zero_grad()
15         loss.backward()
16         optimizer.step()
17
18     b = images.size(0)
19     total_loss += loss.item() * b
20     total_r2 += weighted_r2_score(targets, preds, TARGET_WEIGHTS) * b
21     n += b
22
23 return total_loss / n, total_r2 / n
24

```

## ▼ Training Loop + Metrics + Best Model Saving

```

1 model = OptimizedCNN(num_outputs=len(TARGET_COLUMNS)).to(DEVICE)
2 optimizer = optim.Adam(
3     model.parameters(),
4     lr=LEARNING_RATE,
5     weight_decay=WEIGHT_DECAY
6 )
7
8 train_losses, val_losses = [], []
9 train_r2s, val_r2s = [], []
10
11 best_val_r2 = -1e9
12
13 for epoch in range(NUM_EPOCHS):
14     tr_1, tr_r2 = run_epoch(model, train_loader, optimizer)
15     va_1, va_r2 = run_epoch(model, val_loader)
16
17     train_losses.append(tr_1)
18     val_losses.append(va_1)
19     train_r2s.append(tr_r2)
20     val_r2s.append(va_r2)
21
22     if va_r2 > best_val_r2:
23         best_val_r2 = va_r2
24         torch.save(model.state_dict(), "best_model_weights.pth")
25         joblib.dump(target_scaler, "target_scaler.pkl")
26
27     print(
28         f"Epoch {epoch+1}/{NUM_EPOCHS} | "
29         f"Train R2: {tr_r2:.4f} | Val R2: {va_r2:.4f}"
30     )
31

```

```

Epoch 1/50 | Train R2: -0.5740 | Val R2: -0.0984
Epoch 2/50 | Train R2: 0.0401 | Val R2: 0.0997
Epoch 3/50 | Train R2: 0.0532 | Val R2: 0.1116
Epoch 4/50 | Train R2: 0.0264 | Val R2: 0.1226
Epoch 5/50 | Train R2: 0.0703 | Val R2: 0.1093
Epoch 6/50 | Train R2: 0.0615 | Val R2: 0.0556
Epoch 7/50 | Train R2: 0.0878 | Val R2: 0.1356
Epoch 8/50 | Train R2: 0.0453 | Val R2: 0.1924
Epoch 9/50 | Train R2: 0.1055 | Val R2: 0.1803
Epoch 10/50 | Train R2: 0.0880 | Val R2: 0.1738
Epoch 11/50 | Train R2: 0.0980 | Val R2: 0.0649
Epoch 12/50 | Train R2: 0.1660 | Val R2: 0.1816
Epoch 13/50 | Train R2: 0.1701 | Val R2: 0.1947
Epoch 14/50 | Train R2: 0.1961 | Val R2: 0.1153
Epoch 15/50 | Train R2: 0.1477 | Val R2: -0.3880
Epoch 16/50 | Train R2: 0.1605 | Val R2: -0.0753
Epoch 17/50 | Train R2: 0.0156 | Val R2: 0.2601
Epoch 18/50 | Train R2: 0.1537 | Val R2: -0.0886
Epoch 19/50 | Train R2: 0.2025 | Val R2: -0.0312
Epoch 20/50 | Train R2: 0.1616 | Val R2: -0.3568
Epoch 21/50 | Train R2: 0.1774 | Val R2: 0.2065
Epoch 22/50 | Train R2: 0.1473 | Val R2: -0.0164
Epoch 23/50 | Train R2: 0.0551 | Val R2: 0.0299
Epoch 24/50 | Train R2: 0.1657 | Val R2: 0.2932
Epoch 25/50 | Train R2: 0.2500 | Val R2: 0.1842
Epoch 26/50 | Train R2: 0.1892 | Val R2: 0.2224

```

Epoch 27/50	Train R <sup>2</sup> : 0.2550	Val R <sup>2</sup> : -0.4257
Epoch 28/50	Train R <sup>2</sup> : 0.1907	Val R <sup>2</sup> : 0.1485
Epoch 29/50	Train R <sup>2</sup> : 0.1430	Val R <sup>2</sup> : 0.0554
Epoch 30/50	Train R <sup>2</sup> : 0.2137	Val R <sup>2</sup> : 0.2602
Epoch 31/50	Train R <sup>2</sup> : 0.2197	Val R <sup>2</sup> : 0.0840
Epoch 32/50	Train R <sup>2</sup> : 0.1808	Val R <sup>2</sup> : 0.0855
Epoch 33/50	Train R <sup>2</sup> : 0.2244	Val R <sup>2</sup> : 0.2070
Epoch 34/50	Train R <sup>2</sup> : 0.2928	Val R <sup>2</sup> : 0.0582
Epoch 35/50	Train R <sup>2</sup> : 0.1811	Val R <sup>2</sup> : -0.2494
Epoch 36/50	Train R <sup>2</sup> : 0.2347	Val R <sup>2</sup> : 0.2028
Epoch 37/50	Train R <sup>2</sup> : 0.1022	Val R <sup>2</sup> : 0.2441
Epoch 38/50	Train R <sup>2</sup> : 0.1132	Val R <sup>2</sup> : 0.0808
Epoch 39/50	Train R <sup>2</sup> : 0.2269	Val R <sup>2</sup> : -0.6006
Epoch 40/50	Train R <sup>2</sup> : 0.1876	Val R <sup>2</sup> : -0.4622
Epoch 41/50	Train R <sup>2</sup> : 0.2043	Val R <sup>2</sup> : 0.0961
Epoch 42/50	Train R <sup>2</sup> : 0.2406	Val R <sup>2</sup> : 0.1442
Epoch 43/50	Train R <sup>2</sup> : 0.1684	Val R <sup>2</sup> : -0.0167
Epoch 44/50	Train R <sup>2</sup> : 0.2829	Val R <sup>2</sup> : 0.2773
Epoch 45/50	Train R <sup>2</sup> : 0.2372	Val R <sup>2</sup> : -0.0062
Epoch 46/50	Train R <sup>2</sup> : 0.1654	Val R <sup>2</sup> : -0.0181
Epoch 47/50	Train R <sup>2</sup> : 0.2651	Val R <sup>2</sup> : -0.3135
Epoch 48/50	Train R <sup>2</sup> : 0.3143	Val R <sup>2</sup> : -0.4012
Epoch 49/50	Train R <sup>2</sup> : 0.2001	Val R <sup>2</sup> : 0.3201
Epoch 50/50	Train R <sup>2</sup> : 0.2963	Val R <sup>2</sup> : 0.2955

## ▼ Average Metrics

```

1 print("\n===== FINAL AVERAGES =====")
2 print("Avg Train Loss:", np.mean(train_losses))
3 print("Avg Val Loss:", np.mean(val_losses))
4 print("Avg Train R2:", np.mean(train_r2s))
5 print("Avg Val R2:", np.mean(val_r2s))
6 print("Best Val R2:", best_val_r2)

```

```

===== FINAL AVERAGES =====
Avg Train Loss: 0.15376009347381298
Avg Val Loss: 0.12544329347709815
Avg Train R2: 0.1520189860787308
Avg Val R2: 0.03890925923983255
Best Val R2: 0.320112943649292

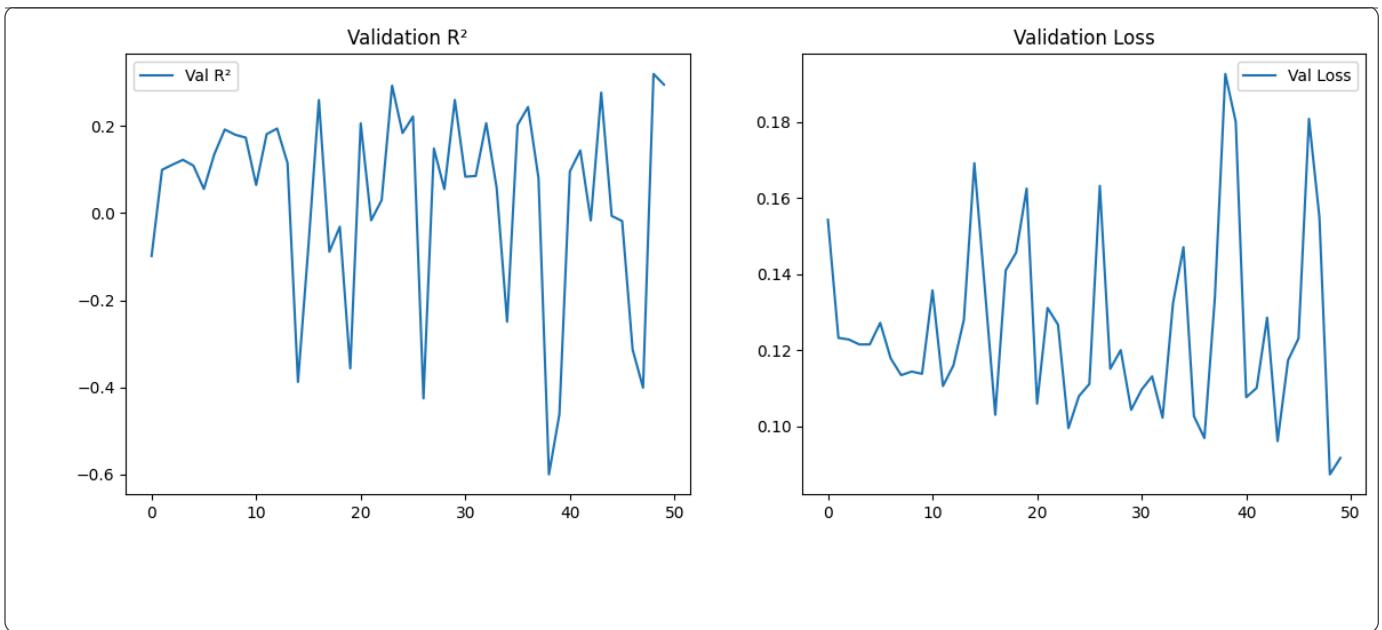
```

## ▼ Plots

```

1 plt.figure(figsize=(14,5))
2 plt.subplot(1,2,1)
3 plt.plot(val_r2s, label="Val R2")
4 plt.title("Validation R2")
5 plt.legend()
6
7 plt.subplot(1,2,2)
8 plt.plot(val_losses, label="Val Loss")
9 plt.title("Validation Loss")
10 plt.legend()
11 plt.show()

```



## ▼ Inference

## ▼ Device & Paths

```

1 DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
2 BASE_PATH = "/kaggle/input/csiro-biomass"
3 TEST_DIR = f"{BASE_PATH}/test"
4 MODEL_PATH = "best_model_weights.pth"
5 SCALER_PATH = "target_scaler.pkl"

```

## ▼ Target Columns (Must Match Training)

```

1 TARGET_COLUMNS = [
2     'Dry_Green_g',
3     'Dry_Dead_g',
4     'Dry_Clover_g',
5     'GDM_g',
6     'Dry_Total_g'
7 ]

```

## ▼ Load Model Weights & Scaler

```

1 best_model = OptimizedCNN(num_outputs=len(TARGET_COLUMNS)).to(DEVICE)
2 best_model.load_state_dict(
3     torch.load(MODEL_PATH, map_location=DEVICE)
4 )
5 best_model.eval()
6
7 target_scaler = joblib.load(SCALER_PATH)
8

```

## ▼ Image Transforms (Validation / Test Only)

```

1 image_transforms = transforms.Compose([
2     transforms.Resize((224, 224)),
3     transforms.ToTensor(),
4     transforms.Normalize(
5         mean=[0.485, 0.456, 0.406],

```

```
6         std=[0.229, 0.224, 0.225]
7     )
8 ])
9
```

## ▼ Test-Time Inference

```
1 submission_rows = []
2
3 with torch.no_grad():
4     for file in tqdm(sorted(os.listdir(TEST_DIR))):
5         image = Image.open(os.path.join(TEST_DIR, file)).convert("RGB")
6         tensor = image_transforms(image).unsqueeze(0).to(DEVICE)
7
8         preds_scaled = best_model(tensor).cpu().numpy()
9         preds = target_scaler.inverse_transform(preds_scaled)[0]
10
11        image_id = file.replace(".jpg", "")
12
13        for name, value in zip(TARGET_COLUMNS, preds):
14            submission_rows.append({
15                "sample_id": f"{image_id}__{name}",
16                "target": float(max(0.0, value)) # clamp negatives
17            })
18
```

100% |██████████| 1/1 [00:00<00:00, 5.59it/s]