

```

1 import os
2 import warnings
3 import numpy as np
4 import pandas as pd
5 import matplotlib.pyplot as plt
6 from PIL import Image
7 from tqdm import tqdm
8
9 import torch
10 import torch.nn as nn
11 import torch.optim as optim
12 from torch.utils.data import Dataset, DataLoader
13 from torchvision import transforms
14
15 from sklearn.model_selection import KFold
16 from sklearn.preprocessing import StandardScaler
17 from xgboost import XGBRegressor
18 from lightgbm import LGBMRegressor
19 from catboost import CatBoostRegressor
20 from sklearn.multioutput import MultiOutputRegressor
21
22 warnings.filterwarnings("ignore")
23
24 # Constants
25 SEED = 42
26 BASE_PATH = "/kaggle/input/csiro-biomass"
27 IMG_SIZE = 224
28 BATCH_SIZE = 16
29 DEVICE = torch.device("cuda" if torch.cuda.is_available() else "cpu")
30
31 np.random.seed(SEED)
32 torch.manual_seed(SEED)
33 if torch.cuda.is_available():
34     torch.cuda.manual_seed_all(SEED)
35
36 TARGET_COLUMNS = ['Dry_Green_g', 'Dry_Dead_g', 'Dry_Clover_g', 'GDM_g', 'Dry_Total_g']
37 TARGET_WEIGHTS = torch.tensor([0.1, 0.1, 0.1, 0.2, 0.5], device=DEVICE)

```

/usr/local/lib/python3.12/dist-packages/sqlalchemy/orm/query.py:195: SyntaxWarning: "is not" with 'tuple' literal. Did you mean if entities is not ():

```

1 # Load and Pivot Data
2 df = pd.read_csv(os.path.join(BASE_PATH, "train.csv"))
3 df = df.pivot_table(index="image_path", columns="target_name", values="target").reset_index()
4 df.fillna(0, inplace=True)
5
6 class BiomassDataset(Dataset):
7     def __init__(self, dataframe, transform=None, scaler=None):
8         self.df = dataframe
9         self.transform = transform
10        self.scaler = scaler
11        # The base path where the 'train' and 'test' folders are located
12        self.base_dir = "/kaggle/input/csiro-biomass"
13
14    def __len__(self):
15        return len(self.df)
16
17    def __getitem__(self, idx):
18        # 1. Get the relative path from the dataframe (e.g., 'train/ID101...jpg')
19        relative_path = self.df.iloc[idx]['image_path']
20
21        # 2. Join it with the correct base directory
22        # This results in: /kaggle/input/csiro-biomass/train/ID101...jpg
23        img_path = os.path.join(self.base_dir, relative_path)
24
25        try:
26            image = Image.open(img_path).convert("RGB")
27        except FileNotFoundError:
28            # Fallback check: if the path structure is slightly different on some OS
29            print(f"Error loading: {img_path}")

```

```

30         raise
31
32     if self.transform:
33         image = self.transform(image)
34
35     targets = self.df.iloc[idx][TARGET_COLUMNS].values.astype(np.float32)
36     if self.scaler:
37         targets = self.scaler.transform(targets.reshape(1, -1)).flatten()
38
39     return image, torch.tensor(targets, dtype=torch.float32)
40
41 # Transforms
42 train_transforms = transforms.Compose([
43     transforms.Resize((IMG_SIZE, IMG_SIZE)),
44     transforms.RandomHorizontalFlip(p=0.5),
45     transforms.RandomRotation(10),
46     transforms.ColorJitter(brightness=0.2, contrast=0.2),
47     transforms.ToTensor(),
48     transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
49 ])
50
51 val_transforms = transforms.Compose([
52     transforms.Resize((IMG_SIZE, IMG_SIZE)),
53     transforms.ToTensor(),
54     transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
55 ])

```

```

1 def weighted_r2_score(y_true, y_pred, weights):
2     mean_y = y_true.mean(0, keepdims=True)
3     ss_res = ((y_true - y_pred) ** 2).sum(0)
4     ss_tot = ((y_true - mean_y) ** 2).sum(0)
5     r2_per_target = 1 - (ss_res / (ss_tot + 1e-8))
6     weighted_r2 = (r2_per_target * weights).sum()
7     return weighted_r2
8
9 def extract_features(model, loader):
10    model.eval()
11    features, targets = [], []
12    with torch.no_grad():
13        for imgs, targs in tqdm(loader, leave=False):
14            imgs = imgs.to(DEVICE)
15            # Standard CLS token extraction
16            output = model(imgs)
17            features.append(output.cpu().numpy())
18            targets.append(targs.numpy())
19    return np.vstack(features), np.vstack(targets)

```

```

1 model_variants = {
2     'ViT-S': 'dinov2_vits14',
3     'ViT-B': 'dinov2_vitb14',
4     'ViT-L': 'dinov2_vitl14'
5 }
6
7 comparison_stats = {}
8
9 for name, repo_key in model_variants.items():
10    print(f"\n--- Testing Variant: {name} ---")
11
12    # Load Backbone
13    backbone = torch.hub.load('facebookresearch/dinov2', repo_key).to(DEVICE)
14    backbone.eval()
15
16    # Prepare Dataloader (no augmentation for pure feature extraction)
17    extract_ds = BiomassDataset(df, val_transforms)
18    extract_loader = DataLoader(extract_ds, batch_size=BATCH_SIZE, shuffle=False)
19
20    # 1. Extract Features
21    X_all, y_all = extract_features(backbone, extract_loader)
22
23    # 2. 5-Fold Cross Validation
24    kf = KFold(n_splits=5, shuffle=True, random_state=SEED)
25    fold_scores = []

```

```

26
27     for fold, (train_idx, val_idx) in enumerate(kf.split(X_all)):
28         X_train, X_val = X_all[train_idx], X_all[val_idx]
29         y_train, y_val = y_all[train_idx], y_all[val_idx]
30
31         # Target Scaling per fold
32         scaler = StandardScaler().fit(y_train)
33         y_train_scaled = scaler.transform(y_train)
34
35         # Train CatBoost
36         cat = MultiOutputRegressor(CatBoostRegressor(
37             iterations=1000,
38             learning_rate=0.05,
39             depth=6,
40             task_type='GPU' if torch.cuda.is_available() else 'CPU',
41             silent=True
42         ))
43
44         cat.fit(X_train, y_train_scaled)
45
46         # Predict and Inverse Scale
47         preds_scaled = cat.predict(X_val)
48         y_pred = np.maximum(scaler.inverse_transform(preds_scaled), 0)
49
50         score = weighted_r2_score(y_val, y_pred, TARGET_WEIGHTS.cpu().numpy())
51         fold_scores.append(score)
52         print(f"Fold {fold+1} R2: {score:.4f}")
53
54         comparison_stats[name] = {
55             'Mean R2': np.mean(fold_scores),
56             'Std Dev': np.std(fold_scores)
57         }
58
59         # Cleanup memory before next model
60         del backbone
61         torch.cuda.empty_cache()
62
63 # 3. Final Summary Table
64 summary_df = pd.DataFrame(comparison_stats).T
65 print("\n==== COMPARISON SUMMARY ===")
66 print(summary_df)

```

```

--- Testing Variant: ViT-S ---
Downloading: "https://github.com/facebookresearch/dinov2/zipball/main" to /root/.cache/torch/hub/main.zip
Downloading: "https://dl.fbaipublicfiles.com/dinov2/dinov2\_vits14/dinov2\_vits14\_pretrain.pth" to /root/.cache/torch/hub/checkpoi
100%|██████████| 84.2M/84.2M [00:00<00:00, 297MB/s]
Fold 1 R2: 0.4629
Fold 2 R2: 0.5751
Fold 3 R2: 0.6487
Fold 4 R2: 0.5853
Fold 5 R2: 0.4793

--- Testing Variant: ViT-B ---
Using cache found in /root/.cache/torch/hub/facebookresearch_dinov2_main
Downloading: "https://dl.fbaipublicfiles.com/dinov2/dinov2\_vitb14/dinov2\_vitb14\_pretrain.pth" to /root/.cache/torch/hub/checkpoi
100%|██████████| 330M/330M [00:01<00:00, 311MB/s]
Fold 1 R2: 0.5424
Fold 2 R2: 0.6086
Fold 3 R2: 0.6294
Fold 4 R2: 0.6024
Fold 5 R2: 0.4638

--- Testing Variant: ViT-L ---
Using cache found in /root/.cache/torch/hub/facebookresearch_dinov2_main
Downloading: "https://dl.fbaipublicfiles.com/dinov2/dinov2\_vitl14/dinov2\_vitl14\_pretrain.pth" to /root/.cache/torch/hub/checkpoi
100%|██████████| 1.13G/1.13G [00:03<00:00, 349MB/s]
Fold 1 R2: 0.5855
Fold 2 R2: 0.5715
Fold 3 R2: 0.6656
Fold 4 R2: 0.5806
Fold 5 R2: 0.3594

==== COMPARISON SUMMARY ====
      Mean R2   Std Dev
ViT-S  0.550276  0.069572
ViT-B  0.569319  0.000178
ViT-L  0.552516  0.102299

```

```
1 plt.figure(figsize=(10, 6))
2 plt.bar(summary_df.index, summary_df['Mean R2'], yerr=summary_df['Std Dev'], capsize=5, color='skyblue')
3 plt.ylabel('Weighted R2 Score')
4 plt.title('DINOv2 Variant Comparison for Biomas s Regression')
5 plt.grid(axis='y', linestyle='--', alpha=0.7)
6 plt.show()
```

