

```

1 import os, torch, numpy as np, pandas as pd
2 from pathlib import Path
3 from PIL import Image
4 from tqdm import tqdm
5 from sklearn.svm import SVR
6 from sklearn.multioutput import MultiOutputRegressor
7 from sklearn.preprocessing import StandardScaler
8 from sklearn.model_selection import KFold, GridSearchCV
9 from sklearn.metrics import make_scorer
10 from transformers import AutoImageProcessor, AutoModel
11 import torchvision.transforms as transforms
12
13 # --- CONFIGURATION ---
14 CFG = {
15     "base_path": Path("/kaggle/input/csiro-biomass"),
16     "model_path": "/kaggle/input/dinov2/pytorch/large/1",
17     "device": torch.device("cuda" if torch.cuda.is_available() else "cpu"),
18     "n_splits": 5,
19     "seed": 42
20 }
21
22 TARGETS = ['Dry_Green_g', 'Dry_Dead_g', 'Dry_Clover_g', 'GDM_g', 'Dry_Total_g']
23 # Official competition weights: Dry_Total_g (0.5), GDM_g (0.2), Others (0.1)
24 COMP_WEIGHTS = np.array([0.1, 0.1, 0.1, 0.2, 0.5])
25 MEAN, STD = [0.485, 0.456, 0.406], [0.229, 0.224, 0.225]

```

2026-01-10 12:47:39.035482: E external/local\_xla/xla/stream\_executor/cuda/cuda\_fft.cc:467] Unable to register cuFFT factory: Attention  
 WARNING: All log messages before absl::InitializeLog() is called are written to STDERR  
 E0000 00:00:1768049259.218180 55 cuda\_dnn.cc:8579] Unable to register cuDNN factory: Attempting to register factory for plu  
 E0000 00:00:1768049259.268875 55 cuda\_blas.cc:1407] Unable to register cuBLAS factory: Attempting to register factory for p  
 W0000 00:00:1768049259.692194 55 computation\_placer.cc:177] computation placer already registered. Please check linkage and  
 W0000 00:00:1768049259.692239 55 computation\_placer.cc:177] computation placer already registered. Please check linkage and  
 W0000 00:00:1768049259.692242 55 computation\_placer.cc:177] computation placer already registered. Please check linkage and  
 W0000 00:00:1768049259.692244 55 computation\_placer.cc:177] computation placer already registered. Please check linkage and

```

1 def weighted_r2_metric(y_true, y_pred):
2     """Calculates Global Weighted R2 based on competition rules."""
3     ss_res = np.sum(COMP_WEIGHTS * np.sum((y_true - y_pred)**2, axis=0))
4     global_mean = np.average(np.mean(y_true, axis=0), weights=COMP_WEIGHTS)
5     ss_tot = np.sum(COMP_WEIGHTS * np.sum((y_true - global_mean)**2, axis=0))
6     return 1 - (ss_res / ss_tot)
7
8 weighted_scoring = make_scoring(weighted_r2_metric, greater_is_better=True)
9
10 def extract_dense_features(df, model):
11     """Extracts patch-based embeddings from DINov2 (Dense Methodology)."""
12     tta_transforms = [
13         transforms.Compose([transforms.ToTensor(), transforms.Resize((224, 224)), transforms.Normalize(MEAN, STD)]),
14         transforms.Compose([transforms.RandomHorizontalFlip(p=1.0), transforms.ToTensor(), transforms.Resize((224,
15         ])]
16     unique_paths = df['image_path'].unique()
17     all_feats = []
18
19     with torch.no_grad():
20         for path in tqdm(unique_paths, desc="GPU Feature Extraction"):
21             img = Image.open(CFG["base_path"] / path).convert("RGB")
22             # Average patch features (ignoring CLS token at index 0)
23             tta_results = [model(aug(img).unsqueeze(0).to(CFG["device"])).last_hidden_state[:, 1:, :].mean(dim=1) for
24             all_feats.append(np.mean(tta_results, axis=0))
25
26     return np.vstack(all_feats), unique_paths

```

```

1 import os, torch, numpy as np, pandas as pd
2 import warnings
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from pathlib import Path
6 from PIL import Image
7 from tqdm import tqdm
8 from lightgbm import LGBMRegressor
9 from xgboost import XGBRegressor

```

```

10 from catboost import CatBoostRegressor
11 from sklearn.multioutput import MultiOutputRegressor
12 from sklearn.preprocessing import StandardScaler
13 from sklearn.model_selection import KFold
14 from sklearn.metrics import r2_score
15 from transformers import AutoModel
16 import torchvision.transforms as transforms
17
18 # --- SILENCE WARNINGS ---
19 warnings.filterwarnings("ignore", category=SyntaxWarning)
20 warnings.filterwarnings("ignore", category=UserWarning)
21 os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'
22
23 # --- CONFIGURATION ---
24 CFG = {
25     "base_path": Path("/kaggle/input/csiro-biomass"),
26     "model_path": "/kaggle/input/dinov2/pytorch/large/1",
27     "device": torch.device("cuda" if torch.cuda.is_available() else "cpu"),
28     "n_splits": 5,
29     "seed": 42
30 }
31
32 TARGETS = ['Dry_Green_g', 'Dry_Dead_g', 'Dry_Clover_g', 'GDM_g', 'Dry_Total_g']
33 COMP_WEIGHTS = np.array([0.1, 0.1, 0.1, 0.2, 0.5])
34 MEAN, STD = [0.485, 0.456, 0.406], [0.229, 0.224, 0.225]
35
36 # --- 1. METRICS ---
37 def weighted_r2_metric(y_true, y_pred):
38     # Calculate R2 for each target
39     r2_scores = []
40     for i in range(y_true.shape[1]):
41         r2_scores.append(r2_score(y_true[:, i], y_pred[:, i]))
42     # Apply competition weights
43     return np.sum(np.array(r2_scores) * COMP_WEIGHTS)
44
45 # --- 2. FEATURE EXTRACTION ---
46 def extract_dense_features(df, model):
47     tta_transforms = [
48         transforms.Compose([transforms.ToTensor(), transforms.Resize((224, 224)), transforms.Normalize(MEAN, STD)])
49         transforms.Compose([transforms.RandomHorizontalFlip(p=1.0), transforms.ToTensor(), transforms.Resize((224
50     )])
51     unique_paths = df['image_path'].unique()
52     all_feats = []
53     with torch.no_grad():
54         for path in tqdm(unique_paths, desc="GPU Feature Extraction"):
55             img = Image.open(CFG["base_path"] / path).convert("RGB")
56             tta_results = [model(aug(img).unsqueeze(0).to(CFG["device"])).last_hidden_state[:, 1:, :].mean(dim=1)
57             all_feats.append(np.mean(tta_results, axis=0))
58
59     feat_matrix = np.vstack(all_feats)
60     feat_cols = [f"feat_{i}" for i in range(feat_matrix.shape[1])]
61     return pd.DataFrame(feat_matrix, columns=feat_cols), unique_paths
62
63 # --- 3. PIPELINE ---
64 def run_pipeline():
65     model_dir = os.path.abspath(CFG["model_path"])
66     model = AutoModel.from_pretrained(model_dir, local_files_only=True, trust_remote_code=True).to(CFG["device"])
67
68     # Data Preparation
69     train_df = pd.read_csv(CFG["base_path"] / "train.csv")
70     train_p = train_df.pivot_table(index="image_path", columns="target_name", values="target").reset_index()
71     X_train_df, _ = extract_dense_features(train_p, model)
72     Y_train_log = np.log1p(train_p[TARGETS].values)
73
74     scaler = StandardScaler()
75     X_train_scaled = pd.DataFrame(scaler.fit_transform(X_train_df), columns=X_train_df.columns)
76
77     # Model Definitions
78     models_to_test = {
79         "LightGBM": MultiOutputRegressor(LGBMRegressor(n_estimators=200, learning_rate=0.05, verbosity=-1, random_
80         "XGBoost": MultiOutputRegressor(XGBRegressor(n_estimators=200, learning_rate=0.05, random_state=CFG["seed"]
81         "CatBoost": MultiOutputRegressor(CatBoostRegressor(n_estimators=200, learning_rate=0.05, verbose=0, rando
82     }

```

```

83
84     cv_results = []
85     kf = KFold(n_splits=CFG["n_splits"], shuffle=True, random_state=CFG["seed"])
86
87     # --- CROSS VALIDATION LOOP ---
88     for name, regressor in models_to_test.items():
89         print(f"\nEvaluating {name}...")
90         fold_scores = []
91
92         for fold, (train_idx, val_idx) in enumerate(kf.split(X_train_scaled, Y_train_log)):
93             X_tr, X_vl = X_train_scaled.iloc[train_idx], X_train_scaled.iloc[val_idx]
94             y_tr, y_vl = Y_train_log[train_idx], Y_train_log[val_idx]
95
96             regressor.fit(X_tr, y_tr)
97             preds = regressor.predict(X_vl)
98
99             # Use competition weighted R2 metric
100            score = weighted_r2_metric(y_vl, preds)
101            fold_scores.append(score)
102            print(f" Fold {fold+1} Weighted R2: {score:.4f}")
103
104            avg_score = np.mean(fold_scores)
105            std_score = np.std(fold_scores)
106            cv_results.append({"Model": name, "Avg_R2": avg_score, "Std_R2": std_score})
107            print(f"Average {name} Weighted R2: {avg_score:.4f} (+/- {std_score:.4f})")
108
109     # --- RESULTS VISUALIZATION ---
110     res_df = pd.DataFrame(cv_results)
111     plt.figure(figsize=(10, 6))
112     sns.barplot(x="Model", y="Avg_R2", data=res_df, palette="viridis")
113     plt.errorbar(x=res_df["Model"], y=res_df["Avg_R2"], yerr=res_df["Std_R2"], fmt='none', c='black', capsize=5)
114     plt.title("Model Comparison: Average Weighted R2 Score (5-Fold CV)")
115     plt.ylabel("Weighted R2 Score")
116     plt.ylim(0, 1.0)
117     plt.savefig("model_comparison.png")
118     print("\nChart saved to model_comparison.png")
119
120     # --- FINAL PREDICTIONS (Using best model) ---
121     best_model_name = res_df.loc[res_df["Avg_R2"].idxmax(), "Model"]
122     print(f"\nFinal Inference using {best_model_name}")
123
124     final_regressor = models_to_test[best_model_name]
125     final_regressor.fit(X_train_scaled, Y_train_log)
126
127     test_df = pd.read_csv(CFG["base_path"] / "test.csv")
128     test_uniq = pd.DataFrame({'image_path': test_df['image_path'].unique()})
129     X_test_df, test_ids = extract_dense_features(test_uniq, model)
130     X_test_scaled = pd.DataFrame(scaler.transform(X_test_df), columns=X_test_df.columns)
131
132     test_preds_log = final_regressor.predict(X_test_scaled)
133     avg_preds = np.maximum(np.expm1(test_preds_log), 0)
134
135     # Submission
136     final_rows = []
137     for i, path in enumerate(test_ids):
138         image_id = Path(path).stem
139         for j, target_name in enumerate(TARGETS):
140             final_rows.append({"sample_id": f"{image_id}_{target_name}", "target": avg_preds[i, j]})
141
142     pd.DataFrame(final_rows).to_csv("submission.csv", index=False)
143     print("Submission saved to submission.csv")
144
145 if __name__ == "__main__":
146     run_pipeline()

```

```
/usr/local/lib/python3.12/dist-packages/sqlalchemy/orm/query.py:195: SyntaxWarning: "is not" with 'tuple' literal. Did you mean
    if entities is not ():
GPU Feature Extraction: 100%|██████████| 357/357 [01:13<00:00,  4.84it/s]
```

```
Evaluating LightGBM...
Fold 1 Weighted R2: 0.5807
Fold 2 Weighted R2: 0.6072
Fold 3 Weighted R2: 0.7106
Fold 4 Weighted R2: 0.6237
Fold 5 Weighted R2: 0.5682
Average LightGBM Weighted R2: 0.6181 (+/- 0.0502)
```

```
Evaluating XGBoost...
Fold 1 Weighted R2: 0.5319
Fold 2 Weighted R2: 0.5610
Fold 3 Weighted R2: 0.6550
Fold 4 Weighted R2: 0.5169
Fold 5 Weighted R2: 0.5046
Average XGBoost Weighted R2: 0.5539 (+/- 0.0539)
```

```
Evaluating CatBoost...
Fold 1 Weighted R2: 0.6191
Fold 2 Weighted R2: 0.6285
Fold 3 Weighted R2: 0.6873
Fold 4 Weighted R2: 0.6492
Fold 5 Weighted R2: 0.5604
Average CatBoost Weighted R2: 0.6289 (+/- 0.0415)
```

Chart saved to model\_comparison.png

```
Final Inference using CatBoost
/tmp/ipykernel_55/1591067300.py:112: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set
```

```
sns.barplot(x="Model", y="Avg_R2", data=res_df, palette="viridis")
GPU Feature Extraction: 100%|██████████| 1/1 [00:00<00:00,  3.41it/s]
Submission saved to submission.csv
```

